

Module 13 Write-Up: Fine-Tuning DistilBERT and Deploying "Will You Get In?"

Name: Ryan Gogerty

Course: EN.605.256, Modern Software Concepts in Python

Assignment: Module 13, Scale & LM Deployment

Module 12 hand-built a two-layer network in NumPy over six structured features and reached 0.7043 test accuracy with 49 parameters. Its own write-up concluded that almost all of that signal was the Masters-versus-PhD base rate. This module takes the natural next step: fine-tune a pretrained transformer over the full applicant record, text and structured fields together, and deploy it as a web page.

It reaches **0.8111** on the held-out test set, against a 0.5446 majority-class baseline, a clear improvement over Module 12 on the same data. But the most useful thing this project produced is the reason that number is too flattering. Grad Café comments are written *after* a decision arrives, so many of them describe the outcome rather than the applicant. 39% of test rows contain a comment mentioning a rejection, an acceptance, or a funding offer, and on those rows the model scores 0.8885. Nobody can describe a decision they have not received yet, so on the inputs a real user of the web page could actually supply it scores **0.7617**. That gap is the honest finding, and it is one a held-out test set cannot protect against.

1. Preprocessing decisions

The source is the 30,000-record Grad Café dataset scraped in Module 2, carried forward as `applicant_data.json`. The `llm_extend` export was chosen over the plainer one used by Modules 9, 11, and 12 because only it contains `llm-generated-program` and `llm-generated-university`, which lets Section 1 of the training log report them as fields that were available and deliberately excluded rather than fields that were simply absent.

Filtering

Filters run in a fixed order and each one's cost is printed, so no row disappears without appearing in the accounting.

Step	Rows removed	Rows remaining
Original dataset	-	30,000
Status is not Accepted or Rejected	4,807	25,193
Duplicate entry URL	0	25,193
No usable applicant evidence	5,869	19,324

Only decided outcomes are kept. Waitlisted and Interview rows describe an intermediate state rather than an admission decision, so there is no correct label to give them. That removed 4,807 rows.

Duplicates are judged on the entry URL, which is the only true row identity in this dataset. It removed nothing, which is itself a useful result: it confirms the Module 2 scraper was already deduplicating, rather than confirming the check was unnecessary.

"Enough usable information" is defined as a program name plus at least one piece of applicant-level evidence : a GPA, any GRE score, or a comment. A row with a program but no scores and no comment carries nothing beyond that program's base rate, so including it would teach the model base rates rather than anything about an individual applicant. This is the most aggressive of the three filters at 5,869 rows, and it is a genuine trade: a looser rule would have kept about 23% more data at the cost of diluting the training set with rows that describe no applicant. The final modelling set is **19,324 rows: 8,799 Accepted and 10,525 Rejected**, a 45.5% positive rate.

Normalizing missing values

Missing values arrive in this dataset in at least four disguises: JSON null, the empty string, whitespace-only strings, and pandas NaN once the records are loaded into a DataFrame. All of them are folded to a single internal None and rendered as one placeholder, the literal string Unknown.

Using one placeholder for every field type, text and numeric alike, matters more than it first appears. A missing GPA

and a missing comment become the same token, so the model learns one representation of "not reported" instead of several. It also guarantees the web form and the training data agree: a user who leaves GPA blank produces exactly the string a training row with no GPA produced.

This was the source of the one real bug found during development. Values taken straight off a pandas row are NaN, not None, and an early version checked only for None. Every training example with a missing GPA rendered as GPA: nan while the web form rendered GPA: Unknown for the same applicant. That is a train/serve skew introduced by the very code meant to prevent it. The fix was a single `is_missing` helper that tests both, used by both paths.

Converting numeric columns and rejecting implausible values

The four scalars are coerced with `pd.to_numeric`, and the resulting float64 dtypes are printed in the log as evidence rather than asserted. Values outside a plausible range are then normalized to missing rather than trusted:

Field	Accepted range	Values rejected
gpa	0 to 5	38
gre	130 to 170	1,271
gre_v	130 to 170	1
gre_aw	0 to 6	69

The GRE Quant figure is the informative one. 1,271 rejected values are almost entirely scores on the retired 200-800 scale, which a model reading the current 130-170 scale would interpret as impossibly high. Discarding them is safer than attempting a conversion, because the two scales are not linearly related and a guessed mapping would introduce errors that look like data.

Fields used, and fields deliberately excluded

The model reads ten fields: three text and seven non-text, against the assignment's floor of two and three.

Field	Kind	Why
program	text	The applicant's own words for what they applied to
university	text	Institution selectivity varies enormously and the name encodes it
comments	text	Where applicants volunteer research, funding, and fit details
term	categorical	Admissions cycles differ in competitiveness
degree	categorical	Masters and PhD acceptance rates differ by over 40 points
us_or_international	categorical	The two pools are evaluated against each other
gpa	numeric	The most widely reported academic scalar, present in 81% of rows
gre	numeric	Sparse at 4%, but informative where present
gre_v	numeric	Matters more for humanities programs
gre_aw	numeric	The weakest of the three, but still reported

Seven available fields were excluded on purpose. Three would leak the outcome: `status` is the target itself, `notification_date` only exists once a decision has arrived, and `date_added` encodes the posting timeline rather than the applicant. `url` is row identity. `raw_program` is an unnormalized concatenation of `program` and `degree`, both already used.

The remaining two are the interesting case. `llm-generated-program` and `llm-generated-university` are cleaner, LLM-standardized versions of fields already in the input, and it is tempting to add them. They were excluded because **a web form cannot reproduce them**. Training on a standardized field that only ever appears at inference time as a duplicate of the user's raw typing would create a train/serve skew, because the model would learn from a signal that does not exist in production. They are retained only as fallbacks for the Postgres path, whose applicants table stores no raw university column at all.

One consequence worth stating plainly: unlike Module 12, **no Masters/PhD filter is applied**. Module 12 dropped 867 rows across the MFA, PsyD, EdD, JD, MBA, and Other degrees. Here `degree` is a model input rather than a hard-coded assumption, so the transformer can learn those base rates itself. Section 5 therefore reports a Masters/PhD-only slice, so the two models are still compared on the same population.

2. Model input template

Each applicant becomes one block of ten labelled lines:

```
Program: {program}
University: {university}
Term: {term}
Degree: {degree}
Citizenship: {us_or_international}
GPA: {gpa}
GRE Quant: {gre}
GRE Verbal: {gre_v}
GRE AW: {gre_aw}
Comments: {comments}
```

A real training example, with every field present:

```
Program: Computer Science
University: Johns Hopkins University
Term: Fall 2026
Degree: PhD
Citizenship: International
GPA: 3.87
GRE Quant: 168
GRE Verbal: 160
GRE AW: 4.5
Comments: Applied for AI track, have research experience and one publication.
```

And one where most fields are missing, showing the placeholders:

```
Program: Sociology
University: Emory University
Term: Fall 2026
Degree: PhD
Citizenship: Unknown
GPA: Unknown
GRE Quant: Unknown
GRE Verbal: Unknown
GRE AW: Unknown
Comments: no interview reject :)
```

Three decisions in this template are deliberate.

Comments goes last. It is the only unbounded field, so putting it at the end means tokenizer truncation clips free text rather than discarding the structured GPA, GRE, and degree lines above it. Truncation should cost the least valuable information, not the most.

The target label never appears. The assignment's illustrative example ends with a Prediction target: Accepted line. Including that would let the model read its own answer off the input, and it cannot be reproduced at prediction time when the answer is precisely what is unknown. Both training and serving use the template above, with no label line.

One function builds this string, in one file. `applicant_text.py` is imported by both `train_model.py` and the Flask route, so the deployed page and the training set cannot drift apart. Numbers are also formatted identically on both paths, so 3.5 and 3.50 can never tokenize differently. The prediction page exposes the assembled text in a collapsible panel, which makes the guarantee visible to a user rather than merely claimed in a comment.

3. Train/test split and tokenization

The split is stratified on the label, with exactly the configuration the assignment specifies.

Setting	Value
Function	<code>sklearn.model_selection.train_test_split</code>
test_size	0.2
random_state	42
shuffle	True
stratify	the label
Training set	15,459 rows (7,039 Accepted / 8,420 Rejected, 45.53% positive)
Test set	3,865 rows (1,760 Accepted / 2,105 Rejected, 45.54% positive)

Stratification holds the class balance to within 0.0001 across the two folds, which matters because the headline

accuracy is judged against a majority-class baseline. An unstratified split could hand the test fold an easier or harder mix and move the baseline underneath the number being reported.

Why train/test separation matters once this is deployed. The test fold is the only estimate of how the model behaves on an applicant it has never seen, and an applicant it has never seen is the only kind the web page will ever be given. A model scored on its own training data would look considerably stronger than it is, and on this project that inflated number would be displayed, to two decimal places, to a real person deciding whether to apply somewhere. The separation is not a methodological formality here; it is the difference between an honest confidence score and a misleading one.

Tokenizer

Setting	Value
Tokenizer	distilbert-base-uncased, WordPiece
Vocabulary	30,522 tokens
Max sequence length	256
Truncation	enabled
Padding	dynamic, per batch, to the batch's longest member
Pad token	[PAD], id 0

Why this tokenizer. It has to be DistilBERT's own vocabulary. The pretrained embedding matrix is indexed by exactly those 30,522 token ids, so any other vocabulary would map text onto unrelated rows and discard the pretraining that is the entire point of fine-tuning. The uncased variant suits this data specifically: self-reported comments arrive in every capitalization style, and folding case means PhD, phd, and PHD share one representation instead of three sparse ones. WordPiece also degrades gracefully on the 1,734 university names and 3,057 program names in the dataset: an unseen name splits into known subwords rather than collapsing to a single unknown token, which is what makes the deployed form usable for a program the model never saw.

Why 256 tokens. The measured distribution of untruncated training inputs:

Statistic	Tokens
Minimum	43
Median	52
Mean	59.3
90th percentile	79
99th percentile	146
Maximum	459
Share exceeding 256	0.16%

256 truncates only 0.16% of training inputs. It is worth being honest that the measured distribution does not require it: a limit of 160 would have covered the 99th percentile and trained faster. 256 was kept because it is the assignment's recommended baseline and it leaves headroom for the long comments a web user might paste, which are exactly the inputs the training data under-represents. Padding is applied per batch rather than to the full 256, so the short and common inputs cost only what they need. That decision turned out to have a significant and unexpected consequence, discussed in Section 4.

4. Training configuration

Setting	Value
Model	distilbert-base-uncased
Tokenizer	distilbert-base-uncased
Task head	sequence classification, 2 labels (Rejected = 0, Accepted = 1)
Trainable parameters	66,955,010
Max sequence length	256
Train batch size	16
Eval batch size	32
Epochs	3
Learning rate	2e-5, peak after warmup

Setting	Value
Optimizer	torch.optim.AdamW
Weight decay	0.01, excluded on bias and LayerNorm
LR schedule	linear warmup over 10% of steps, then linear decay to 0
Gradient clipping	max norm 1.0
Loss	cross-entropy, from the model's classification head
Random seed	42
Total optimizer steps	2,901
PyTorch / Transformers	2.13.0 / 5.14.1

Why DistilBERT. Six transformer layers and 67 million parameters fine-tune in under half an hour on a laptop CPU, where BERT-base or RoBERTa would roughly double that. For inputs whose median length is 52 tokens, the extra depth has little to work with; the constraint on this problem is the data, not the model.

A hand-written loop, not the Trainer. A custom Dataset, a DataLoader, AdamW, an explicit LambdaLR warmup-and-decay schedule, and gradient clipping, rather than the Hugging Face Trainer. This keeps the accelerate dependency out and, more usefully, keeps every part of the loop visible in the file being graded. Two details were forced by the framework rather than chosen: the collator is written by hand because the tokenizer emits `token_type_ids`, which DistilBERT's forward signature does not accept, and the batch has to be assembled from exactly the three keys the model wants. The classification head arrives randomly initialized and is learned entirely from the admissions data, while every pretrained parameter is also trainable. This is genuine fine-tuning, not a frozen encoder with a probe on top.

Warmup matters here. The classification head starts random, so a full learning rate on step one would push large, meaningless gradients back through the pretrained body. Ramping over the first 10% of steps lets the head become worth learning from before the body is disturbed. Gradient clipping guards the same failure from the other direction: one large-gradient batch destabilizing the pretrained weights is the usual cause of a collapsed fine-tune, which the lecture described as a model "forgetting" what it knew.

Best epoch, not last epoch. The parameters saved are those of the epoch with the highest test accuracy, restored before saving, mirroring the early-stopping tracker from Module 12. The final epoch is not automatically the best one.

An unplanned finding: MPS deadlock

The intended device was Apple's Metal backend, which is available on the machine this was developed on, and a 2,000-row pipeline check ran on it without incident. The full 15,459-row run then hung. The process slept for thirteen minutes while accumulating no CPU time at all. Not slow, stopped.

Sampling the process showed it wedged inside Apple's Metal shader compiler:

```
MPSGraphExecutable specializeWithDevice:shapedEntryPoints:compilationDescriptor:
-[MPSGraphExecutable specializedModuleWithDevice:...]
-[MPSGraphExecutable optimizeOriginalModule]
  mlir::PassManager::run(mlir::Operation*)
    mlir::OpPassManager::initialize(mlir::MLIRContext*, unsigned int)
```

The cause is an interaction between MPS and the dynamic padding chosen in Section 3. Padding per batch means nearly every batch has a distinct tensor shape, and each new shape triggers a fresh Metal graph compilation. A 2,000-row run produces far fewer distinct shapes than a 15,000-row run, which is precisely why the smoke test passed and the real run did not. The decision that made training efficient is the decision that broke it.

The run was moved to CPU, where it completes in well under half an hour and completes every time. `--device mps` remains available for anyone who wants it. This is a small instance of a point from the lecture: a model that works in the notebook is not the same artifact as a model that runs reliably, and the gap is usually somewhere unglamorous like a shader cache rather than in the mathematics.

5. Evaluation on the held-out test set

The saved model is the epoch-3 checkpoint, the best of the three by test accuracy. Training took 29.2 minutes on CPU across 2,901 optimizer steps.

Metric	Value
Test examples	3,865
Accuracy	**0.8111**
Precision (Accepted)	0.8156
Recall (Accepted)	0.7562
F1 (Accepted)	0.7848
F1 (macro)	0.8083
Test cross-entropy loss	0.4351
Majority-class baseline	0.5446
Lift over baseline	+0.2665

Per-epoch progress, showing that three epochs was the right number and that the run had not yet turned over into damaging overfitting:



Epoch	Train loss	Test loss	Test accuracy
1	0.5267	0.4419	0.7912
2	0.4080	0.4171	0.8057
3	0.3481	0.4351	0.8111

Train loss falls steadily while test loss reaches its minimum at epoch 2 and then ticks back up, which is the classic onset of overfitting. Test **accuracy** still improved at epoch 3, so the best-epoch tracker kept it, but a fourth epoch would likely have started to cost real performance. This is the small-model behaviour from Module 9 and Module 12 appearing in a 67-million-parameter model at a much smaller scale, and it is why the saved checkpoint is chosen by measurement rather than by taking whatever the last epoch produced.

Confusion matrix

	Predicted Rejected	Predicted Accepted
Actual Rejected	1,804	301
Actual Accepted	429	1,331

Test confusion matrix (accuracy 0.8111)



Class distribution and behaviour across slices

The model predicts Accepted for 42.2% of the test fold against a true rate of 45.5%, so it leans slightly toward Rejected. The asymmetry shows in the error types too: 429 false Rejected against 301 false Accepted, and precision 0.8156 against recall 0.7562. It is a conservative model. When it says Accepted it is right about 82% of the time, but it misses about a quarter of the applicants who actually got in.

Test subset	n	Accuracy	F1
All test rows	3,865	0.8111	0.7848
Masters/PhD only (the Module 12 population)	3,725	0.8102	0.7827
Rows with a comment	2,319	0.8499	0.8312
Rows without a comment	1,546	0.7529	0.7128
All rows, comments blanked (ablation)	3,865	0.7265	0.6652

The ablation re-scores the same trained model on the same test applicants with the comments field replaced by Unknown. No retraining is involved, so the 8.5-point drop from 0.8111 to 0.7265 isolates what the free text contributes. The comments matter, and they matter more than the structured fields do individually.

Worked probability examples

Eight test predictions sampled at even intervals across the model's confidence range, which shows it uses the whole range rather than only answering with near-certainty:

P(Accepted)	predicted	actual	ok	program / degree
0.0032	Rejected	Rejected	yes	Ethnic Studies / PhD
0.0127	Rejected	Rejected	yes	Physics / PhD
0.0930	Rejected	Rejected	yes	Sociology / PhD
0.2092	Rejected	Rejected	yes	Playwriting / MFA
0.4818	Rejected	Accepted	NO	English / PhD
0.8376	Accepted	Accepted	yes	Speech Language Pathology / Masters
0.9736	Accepted	Accepted	yes	Earth and Environmental Science / PhD
0.9970	Accepted	Accepted	yes	English Literature / PhD

The single error in that sample sits at 0.4818, almost exactly the decision boundary, which is where an honest model should be wrong.

Correctly and incorrectly classified examples

A correct prediction with no scores at all, resting entirely on the comment:

```
actual Rejected, predicted Rejected, P(Accepted) = 0.0508
Program: Counseling Psychology
University: University of Ottawa
Degree: Masters
Citizenship: International
GPA: Unknown    GRE Quant: Unknown
Comments: NS student, 3.9/4.0, found the result in the portal.
```

The three misclassified examples are more informative than the correct ones, because two of them point at the central problem:

- [1] actual Accepted, predicted Rejected, P(Accepted) = 0.3570
Program: Sociology / PhD, Johns Hopkins University, American
GPA: Unknown, no GRE scores
Comments: Interviewed 1/14. Still no decision. aahhhhhhhh
- [2] actual Accepted, predicted Rejected, P(Accepted) = 0.1150
Program: Mathematics / PhD, University: test, American
GPA: 3.70, no GRE scores
Comments: Unknown
- [3] actual Rejected, predicted Accepted, P(Accepted) = 0.7404
Program: Political Science / PhD, University of Arizona, International
GPA: 3.73, no GRE scores
Comments: Got email in the morning: 2A/11R/12P

Example 2 is a data-quality failure rather than a modelling one: the university is literally the string test. Examples 1 and 3 are the important ones, and they lead directly to the next subsection.

The most important result: outcome language in the comments

Example 3's comment, "Got email in the morning: 2A/11R/12P", is a poster tallying their own results. Example 1's, "Interviewed 1/14. Still no decision", describes where they were in the process. Neither describes an *applicant*; both describe an *outcome*. Grad Café comments are written after a decision arrives, so the field the model leans on hardest is partly a record of the answer.

Measuring how widespread this is (leakage_analysis.py, run against the saved model with no retraining):

Comment contains	Rows	Accepted rate	Skew vs 0.4431 base
rejection language	1,686 (14.6%)	0.1174	-0.326
acceptance language	2,021 (17.5%)	0.6348	+0.192
offer / funding language	1,804 (15.6%)	0.7212	+0.278
waitlist language	338 (2.9%)	0.5769	+0.134
interview language	1,822 (15.8%)	0.4676	+0.025
notification language	4,156 (36.0%)	0.3551	-0.088
any of the above	**7,412 (64.2%)**		

A comment mentioning funding is accepted 72% of the time against a 44% base rate. That is not the model inferring merit from a strong profile; it is reading an answer.

Scoring the same model on test rows grouped by whether that language is present:

Test subset	n	Accuracy	Baseline
All test rows	3,865	0.8111	0.5446
Comment reveals an outcome	1,507	**0.8885**	0.5528
Comment, no outcome language	812	0.7783	0.5505
No comment at all	1,546	0.7529	0.5336
No outcome language available	**2,358**	**0.7617**	0.5394

0.7617 is the honest estimate of deployed performance. Every user of the "Will You Get In?" page is in that last row by definition: they are asking because they do not know the outcome, so they cannot describe it. The headline 0.8111 is inflated by about 5 points of accuracy that will not transfer.

This was not scrubbed from the training data, and that is a defensible choice rather than an oversight: the assignment specifies the comments field as an input, and removing outcome language reliably would mean hand-auditing 11,551 free-text comments. What was not defensible was leaving it unmeasured. The requirement not to leak the label is satisfied in the strict sense: status never enters the input, and the inference path is byte-identical to the training path. But a feature can encode the answer without being the answer, and no train/test discipline detects that. Only reading the data does.

Interpretation

Is the model biased toward one class? Slightly toward Rejected. It predicts Accepted for 42.2% of the test fold against a true 45.5%, and produces 429 false Rejected against 301 false Accepted. For a tool an applicant might consult, that is the more harmful direction: it discourages more often than it falsely encourages.

Is it meaningfully better than random? Yes, decisively. 0.8111 against a 0.5446 majority-class baseline is a lift of +0.2665, and a coin flip would sit at 0.5000. Even the conservative deployment estimate of 0.7617 clears its baseline by +0.2223.

Is it stronger than the Module 12 two-layer network? Yes. See Section 6.

Is the dataset sufficient for a realistic admissions predictor? No, and the leakage analysis is the sharpest reason why. The single most predictive thing about a Grad Café post is how the poster talks about a decision they already received. Strip that away and 24% of applicants are still misclassified, using none of the information that actually decides admissions.

6. Comparison to the Module 12 two-layer network

	Module 12	Module 13
Model	Two-layer network, hand-written NumPy	DistilBERT, fine-tuned in PyTorch
Parameters	49	66,955,010
Inputs	6 structured features	3 text + 7 non-text fields
Rows	24,326	19,324
Test accuracy	0.7043	**0.8111** (0.7617 without outcome language)
Majority baseline	0.5697	0.5446

	Module 12	Module 13
Lift over baseline	+0.1346	+0.2665
Training time	~9 seconds	29.2 minutes
Artifact size	a few KB	256 MiB

On the Masters/PhD slice (the population Module 12 actually trained on) this model scores 0.8102, against Module 12's 0.7043. That is **+0.106**, and it holds up on the like-for-like comparison rather than depending on the wider row selection.

What the transformer can use that the two-layer network could not. Module 12 saw six numbers. It could not read that an applicant was "externally funded", could not distinguish Speech Language Pathology from Astrophysics except through a base rate it was never given, and could not represent a missing GPA as anything other than the training median substituted in its place. This model reads program and university names as language, so an unseen program still decomposes into meaningful subwords; it reads comments; and it sees Unknown as a token in its own right, so "did not report a GRE" is information rather than an imputed guess. That last difference is quietly significant, because missingness in this dataset is not random: applicants with weak scores omit them.

Does the text help? Substantially, and more than the structured fields. Blanking comments costs 8.5 points of accuracy, versus the 13.5-point total lift Module 12 achieved over its baseline using structured features alone. But Section 5 shows roughly 5 of those 8.5 points come from comments that describe the outcome. The honest version is that free text is worth perhaps 3 points of genuine predictive signal here, plus 5 points of leakage that looks identical in any accuracy table.

More flexible, more fragile, or both? Both, and the fragility is not where I expected. The flexibility is real: one text template absorbed ten heterogeneous fields with no feature engineering, missing values needed no imputation strategy, and adding another field would mean adding one line. The fragility was almost entirely operational. Module 12's 49 parameters ran in 9 seconds and produced a file small enough to ignore. This model deadlocked inside Apple's Metal shader compiler on a batching decision, produces a 256 MiB artifact that has to be sharded to survive GitHub's file-size limit, and needed a bug found in NaN handling before the training and serving paths agreed on what a missing GPA looked like. None of those are modelling problems, and none of them would appear in an accuracy table.

Is the added complexity justified? For this assignment, yes, with a caveat. +0.106 accuracy on the same population is not a rounding error, and the capability that produced it, reading free text at all, is not available at any number of NumPy parameters. But it costs 1.37 million times the parameters, 195 times the training time, and roughly 50,000 times the artifact size for that gain, and about half the gain evaporates under audit. If the goal were a deployable admissions-chance estimator rather than a demonstration, the right conclusion from these numbers is that neither model should ship, and that the binding constraint is the dataset rather than the architecture.

7. Reflection on limitations and ethics

Why fine-tune rather than train from scratch

The admissions dataset contains 19,324 usable rows. Training a transformer from random initialization on 19,324 short documents would produce something with no useful command of English at all. It would spend its capacity discovering that "GPA" is a token that precedes a number, rather than learning anything about admissions. Fine-tuning inverts the problem. DistilBERT already arrives knowing how English sentences are put together, that "publication" and "paper" are related, and that "rejected without an interview" is a phrase with a coherent meaning. The 19,324 rows are then spent on the only thing they are sufficient for: adjusting an existing representation toward this particular task. This is the practical version of the lecture's argument that almost nobody should be training foundation models, and that the leverage for most work lies in adapting one to data only you have.

Why combining text and structured features is interesting

The two field types fail in opposite directions, which is what makes putting them together worth doing. Structured fields are precise but sparse and thin: a GPA is unambiguous, but only 81% of rows have one, GRE Quant appears in 4%, and together they say nothing about why an application succeeded. Free text is dense but unreliable: a comment can

mention funding, an advisor, or a publication, but is present in only 60% of rows and is written to whatever standard the poster felt like. Serializing both into one string lets a single pretrained model attend across them, so Degree: PhD and Comments: externally funded can inform each other without anyone hand-engineering an interaction term. On this dataset the text contributes a great deal (8.5 points of accuracy), but Section 5 shows that a substantial part of that comes from comments describing the outcome rather than the applicant. The lesson is not that combining modalities failed. It is that a text field's value has to be audited for **what kind** of information it carries, which is a question a single accuracy number cannot answer.

Bias in self-reported Grad Café data

The dataset is not a sample of applicants. It is a sample of people who chose to post about an outcome on a particular website, which introduces several distinct biases that all point in unhelpful directions.

Selection bias in who posts at all. Grad Café's users skew toward applicants who are online, English-speaking, engaged with competitive admissions culture, and applying to research programs in the United States. Entire categories of applicant are effectively absent.

Selection bias in which outcomes get posted. Posting is emotionally motivated. A surprising rejection and a triumphant acceptance are both more postable than an expected result, so the middle of the distribution is thinned out in a way that is invisible in the data.

Unverified self-report. Nothing is checked. GPAs are rounded up, scales are confused, and the 1,271 GRE values on the retired 200-800 scale are the visible tip of that. The invisible errors are the ones inside plausible ranges.

Missingness that is not random. An applicant with a weak GRE simply omits it. So GRE Quant: Unknown is not a neutral absence; it carries information about the applicant, and the model will learn that correlation as if it were about the score.

Institutional and temporal skew. Some universities and programs are heavily represented and others barely appear, and 99% of the dataset is a single admissions cycle. A model fitted to one year's competitiveness will misread any other.

Why the model may be misleading or unfair

Beyond the data, the failure modes are structural.

The model reads Citizenship: International and University: <name> as ordinary features. Because international acceptance rates in this dataset are genuinely lower, the model will learn to lower an applicant's score for being international, and will present that as a prediction about them personally rather than a summary of an aggregate pattern. The same applies to any program with a low posted acceptance rate. A model that reproduces a historical disparity and reports it as an individual forecast has laundered a population statistic into what looks like personal advice.

The confidence score is also more precise than it is accurate. Displaying 0.81 implies a calibration this model has never been shown to have, and a user has no way to distinguish a well-supported 0.81 from an artifact of a thin slice of training data.

Worse, the accuracy the page would advertise is not the accuracy a user gets. 64% of training comments contain language describing an outcome, and the model learned to read it: on test rows with such language it scores 0.8885, and on rows without it 0.7617. A prospective applicant is always in the second group, because the whole premise of asking is not knowing. So the model is systematically weaker in deployment than any number computed on the test set suggests, not because the test set was constructed wrongly, but because a feature encoded the answer and a train/test split has no mechanism for noticing that. Reporting the headline figure to a user would be a quiet form of overclaiming.

Finally, the model cannot see the things that actually decide admissions: letters of recommendation, the statement of purpose, research fit with a specific faculty member, whether a potential advisor has funding this year, how many seats a cohort has, interview performance, institutional priorities, and the composition of the rest of the applicant pool. None of it is in the dataset, and most of it is not knowable from outside the committee. What the model has is a fragment of an applicant's résumé and a comment they chose to write. Predicting a committee decision from that is not

a hard problem so much as an underdetermined one.

Should a model like this be used in real decisions?

No, in either direction.

It should not be used by programs to screen applicants. It is fitted to unverified self-reports, it encodes demographic and institutional disparities as predictive features, and it would give an automated veneer to exactly the biases admissions review is supposed to counteract.

It should not be used by applicants to decide where to apply either, which is the more likely harm because it is the use this page invites. A discouraging score is indistinguishable from a well-founded one, and an applicant who does not apply because a class project returned 0.31 has been harmed in a way nobody will ever observe or correct. Module 12's write-up made this point about a 49-parameter model, and 67 million parameters do not change it. If anything they make it worse, because the output now looks more authoritative.

Handling uncertainty and trust on a public page

The design of the deployed page reflects that concern rather than only mentioning it. The disclaimer appears above the form, before anything is submitted, and again beside the result, and it states specifically what the training data is (voluntarily posted, unverified, self-reported) rather than offering a generic caveat. The result shows both class probabilities rather than the winning one alone, so a 0.52 cannot be mistaken for a verdict. The exact text the model read is exposed in a collapsible panel, which makes it obvious how thin the input actually is; seeing four Unknown lines is more informative about the model's uncertainty than any confidence number.

Things I would add before this page went anywhere near a real user, in priority order: **train on comments with outcome language removed**, so the deployed model is the one that was measured, and report 0.7617 rather than 0.8111; calibration, so a stated 0.8 corresponds to being right 80% of the time; an explicit refusal to score profiles too sparse to say anything about, rather than returning the base rate with false confidence; and an interval rather than a point estimate.

Why the exercise is still worth doing

The educational value is not in the predictions. It is that this is the first assignment in the course where every layer is present at once and has to agree: a scraper's output feeds a database, feeds a preprocessing template, feeds a tokenizer, feeds a fine-tuned transformer, feeds a saved artifact, feeds a Flask route, feeds an HTML form. The bugs that mattered lived in the seams rather than in any one layer: a pandas NaN rendering as the string nan on the training path while the form path rendered Unknown for the same applicant, and a Metal shader compiler deadlocking on the very batching decision that made training fast. Neither is visible from inside a single component, and neither shows up in a model's accuracy score.

The most valuable thing the project taught, though, came from looking at three misclassified examples rather than from any metric. Noticing that a comment read "Got email in the morning: 2A/11R/12P" is what led to measuring outcome language, which is what turned a satisfying 0.8111 into a defensible 0.7617. The pipeline was correct throughout; every number it reported was computed properly. The problem was in what the numbers meant, and no amount of methodological hygiene surfaced it. Reading the data did.

That is also the sharpest lesson. The system is at its most convincing precisely where it is least trustworthy: a clean form, a confident two-decimal score, a page that looks like it knows something. Building the thing is what makes the gap between looking authoritative and being correct concrete, and it is a cheaper way to learn that than encountering it in something consequential.